

# VoxNeuro: Leakage-Safe Parkinson’s Screening Support via Multi-View Grassmannian Speech Analysis

Laiba Khan<sup>1\*</sup> and Manas Waghe<sup>2</sup>

<sup>1</sup> Lexington High School, Lexington, MA 02421, USA; 27stu469@lexingtonma.org

<sup>2</sup> Worcester Polytechnic Institute, Worcester, MA 01609, USA; mywaghe@wpi.edu

\* Correspondence: 27stu469@lexingtonma.org

**Abstract:** Parkinson’s disease (PD) can alter phonation and articulation, making repeated speech a low-burden digital biomarker candidate. Public speech datasets contain several recordings per participant, and record-wise cross-validation can therefore place one individual in both training and test sets, so reported scores partly measure recognition of an observed person rather than generalization to an unseen participant. VoxNeuro treats the subject, not the recording, as the statistical observation. Each subject’s three recordings, standardized within the outer training fold, define a rank-three Grassmann subspace and a Euclidean mean–variability vector. Gaussian kernels on the two views are fused into a positive-semidefinite kernel for support vector classification. For imbalanced data, minority subjects are synthesized along Grassmann geodesics with matched Euclidean interpolation. Under five-fold subject-disjoint evaluation, the fused model attained the highest mean balanced accuracy and macro-F1 among the included comparators: 0.850 and 0.847 on UCI-489 and, on PD-252,  $0.739 \pm 0.012$  and  $0.752 \pm 0.013$  (mean  $\pm$  standard deviation, 24 augmentation seeds); margins over a matched Euclidean ablation were positive in 23/24 (macro-F1) and 16/24 (balanced-accuracy) draws. Exploratory permutation analysis identified the tunable Q-factor wavelet transform block as the largest Euclidean-summary contributor. These results support subject-disjoint evaluation and motivate geometry-aware representations for repeated-measures speech biomarkers.

**Keywords:** Parkinson’s disease; speech biomarker; Grassmann manifold; kernel fusion; imbalanced learning; subject-disjoint cross-validation; support vector machine; digital health

## 1. Introduction

Parkinson’s disease (PD) is a progressive neurodegenerative disorder involving nigrostriatal dopaminergic loss, basal ganglia circuit dysfunction, and heterogeneous motor and non-motor manifestations [1–3]. Its diagnosis remains clinical [4,5], and its global burden has grown faster than that of any other neurological disorder [6,7]. Disordered speech is among the most common motor consequences: dysarthria, hypophonia, and related vocal-tract dysfunctions affect the large majority of patients over the disease course [8,9]. Quantitative acoustic analysis detects these changes even in early and untreated disease [10]. Nonlinear dysphonia measures separate disordered from healthy voices [11], and repeated noninvasive speech tests have supported research on low-burden remote monitoring and progression tracking [12–14]. Automated speech analysis has revealed prodromal changes in at-risk cohorts [15]. Machine-learning approaches to voice-based PD detection have accordingly expanded rapidly, as documented by recent systematic reviews [16,17]. A digital biomarker, however, is clinically credible only when its evaluation matches its intended unit of use: the patient.

Public PD speech datasets often contain multiple recordings per participant. When recordings are shuffled independently, subject-specific anatomical and acquisition signatures can occur on both sides of a record-wise split, and the evaluation then estimates recognition of another recording from a previously observed person rather than generalization to an unseen person. This failure mode is a specific instance of data leakage, a recognized driver of over-optimistic results across applied machine learning [18,19]. This effect has been demonstrated directly in clinical prediction: record-wise cross-validation can inflate accuracy substantially relative to subject-wise evaluation of the very same pipeline [20], a finding that prompted an extended methodological discussion [21]. Subject characteristics can also drive apparent performance even when splits look innocuous [22]. Related biases arise whenever model selection and error estimation share data [23,24], and cross-validation on small samples carries


wide error bars even under favorable conditions [25–27]. The UCI-489 corpus analyzed here contains three dependent replications from each of 80 subjects [28]; replication-aware analysis has been proposed for exactly this cohort [29].

Methodologically, three research strands converge in this study. First, engineered speech descriptors: dysphonia, cepstral, wavelet, and tunable Q-factor wavelet transform (TQWT) descriptor families have been used to characterize PD-related voice changes [12,30,31]. Second, subspace geometry: collections of related observations can be represented as points on the Grassmann manifold  $\text{Gr}(r, d)$ , whose quotient structure removes arbitrary basis choices [32–34]; this representation underpins subspace-based learning in computer vision and pattern recognition [35,36], and kernels on such manifolds require care because Gaussian functions of arbitrary manifold distances are not automatically positive semidefinite (PSD) [37]. Third, kernel fusion: reproducing-kernel Hilbert space (RKHS) theory [38–40] and multiple-kernel learning [41–43] combine heterogeneous similarities while preserving convex support vector optimization. Class imbalance, finally, is routinely addressed by synthetic oversampling [44–46], whose behavior in high-dimensional feature spaces warrants caution [47]. This study combines these strands under an evaluation protocol that treats the patient as the statistical unit.

Deep architectures, including convolutional, recurrent, and transformer-based models, are increasingly applied to voice-based PD detection and scale well when large corpora are available [16]. Smartphone-scale collection efforts such as mPower point in that direction [48,49], and technology roadmaps for PD anticipate exactly this kind of remote, high-frequency assessment [50–52]. The cohorts relevant to replicated-phonation screening, however, typically contain tens to a few hundred subjects. In this regime, the present study deliberately combines structured geometric priors with a convex, auditable classical learner instead of high-capacity end-to-end models: each component (representation, kernel, solver, and resampling step) has an explicit mathematical description that can be verified directly against its implementation.

VoxNeuro aligns the mathematical unit of analysis with the clinical unit of interest: each subject is represented by the complete set of repeated measurements. A fixed-rank representation based on singular value decomposition (SVD) encodes the joint orientation of the repetitions, while a Euclidean summary preserves feature-space location and within-subject variability. A positive-semidefinite direct-sum kernel combines the two views, and evaluation is subject-disjoint throughout (Figure 1). For imbalanced data, synthetic subjects are generated along manifold geodesics rather than by interpolating individual recording rows. The study had two objectives: (i) to formalize why record-wise and subject-disjoint evaluation estimate different risks; and (ii) to assess whether adding the Grassmann kernel to a matched nonlinear Euclidean model improves subject-level classification.

The principal contributions are as follows: (i) a formal subject-disjoint risk definition together with a quantitative account of the optimistic bias induced by record-wise splitting; (ii) a basis-invariant Grassmann representation of repeated speech; (iii) provably positive-semidefinite fusion of geometric and Euclidean kernels; (iv) geodesic minority oversampling coupled to the Euclidean view; and (v) a separate research-stage web application for guided speech capture and longitudinal result visualization, not used in the present analyses. Evaluation across datasets uses leakage-safe, dataset-appropriate comparators, and on PD-252 a matched nonlinear Euclidean ablation with a radial basis function (RBF) kernel isolates the contribution of the Grassmann kernel term.



**Figure 1.** Schematic overview of the VoxNeuro framework. (A) Each subject contributes three repeated recordings and one diagnostic label; the recordings are standardized with statistics estimated inside the outer training fold only and stacked into a matrix  $M_i$ . (B) A thin singular value decomposition (SVD) yields the Grassmann view  $[Q_i]$ , while coordinate-wise mean and dispersion yield the Euclidean view  $s_i$ . (C) Gaussian kernels on the two views are combined into a positive-semidefinite (PSD) fused kernel feeding a class-weighted support vector machine (SVM). (D) Evaluation is five-fold subject-disjoint cross-validation (CV); Grassmann-geodesic synthetic minority oversampling (G-SMOTE) operates inside training folds only, and fold scores are summarized as mean  $\pm$  standard error (SE). PD, Parkinson’s disease.

## 2. Materials and Methods

### 2.1. Datasets

Two public, de-identified cohorts were analyzed (Table 1). The first, denoted UCI-489, comprises 80 subjects (40 PD, 40 control) with three recordings of sustained /a/ per subject, each recording described by 44 replicated acoustic descriptors and one dataset-provided binary sex/gender variable [28,29]. The second, denoted PD-252 [53], comprises 252 subjects (188 PD, 64 control) with three recordings per subject, described by 752 descriptors spanning baseline acoustic, vocal/time-frequency, Mel-frequency cepstral coefficient (MFCC) [54], wavelet, TQWT [55], and intrinsic mode function/empirical mode decomposition (IMF/EMD) [56] families [30,31]. It also includes one dataset-provided binary sex/gender variable, which constitutes the demographic family shown in Figure 8. The coded variable was included as a fixed model input and was never used to construct folds. This study was a secondary computational analysis of public data and involved no participant recruitment.

**Table 1.** Dataset regimes and evaluation design.

| Attribute              | UCI-489       | PD-252                           |
|------------------------|---------------|----------------------------------|
| Subjects               | 80            | 252                              |
| PD / control           | 40 / 40       | 188 / 64                         |
| Recordings per subject | 3             | 3                                |
| Modeled variables      | 45            | 753                              |
| Acoustic regime        | Sustained /a/ | Multi-family engineered features |
| Oversampling           | Inactive      | Minority control class only      |

## 2.2. Subject-Disjoint Problem Formulation

Let  $R_i = \{(x_{ij}, y_i) : j = 1, \dots, m_i\}$  denote all recordings from subject  $i$ ; Table 2 summarizes the notation used throughout. For outer fold  $k$ , let  $I_k^{\text{tr}}$  and  $I_k^{\text{te}}$  be the training and held-out subject-index sets. Leakage-safe evaluation requires

$$I_k^{\text{tr}} \cap I_k^{\text{te}} = \emptyset, I_k^{\text{tr}} \cup I_k^{\text{te}} = \{1, \dots, n\}. \quad (1)$$


All recordings in  $R_i$  are assigned to the same partition as subject  $i$  (Figure 2A,B). For a pointwise classification loss  $\ell$ , the fold risk is evaluated once per held-out subject,

$$\hat{R}_k(f) = \frac{1}{|I_k^{\text{te}}|} \sum_{i \in I_k^{\text{te}}} \ell(f(Z_i), y_i), \quad (2)$$

where  $Z_i$  is the complete subject-level representation. Equations (1) and (2) define the intended estimand, generalization to an unseen person rather than to an unseen row from a known person; the reported balanced accuracy and macro-F1 are fold-level aggregate metrics defined in Equation (15).

**Table 2.** Principal notation.

| Symbol                                                  | Meaning                                                                    |
|---------------------------------------------------------|----------------------------------------------------------------------------|
| $x_{ij} \in \mathbb{R}^d$                               | feature vector of recording $j$ of subject $i$                             |
| $R_i, y_i$                                              | all recordings of subject $i$ and its diagnostic label                     |
| $I_k^{\text{tr}}, I_k^{\text{te}}$                      | training and held-out subject-index sets of outer fold $k$                 |
| $M_i$                                                   | standardized $d \times 3$ recording matrix of subject $i$                  |
| $Q_i, [Q_i]$                                            | orthonormal basis from the thin SVD and its subspace class                 |
| $\text{Gr}(r, d), \text{St}(r, d)$                      | Grassmann and Stiefel manifolds of rank $r$ in $\mathbb{R}^d$              |
| $P_i = Q_i Q_i^\top$                                    | orthogonal projector representing $[Q_i]$                                  |
| $\theta_{ij, q'}, d_{\text{ch}}$                        | principal angles and chordal distance between subspaces                    |
| $s_i \in \mathbb{R}^{2d}$                               | Euclidean mean–dispersion summary of subject $i$                           |
| $K_G, K_E, K; w$                                        | Grassmann, Euclidean, and fused kernels; fusion weight                     |
| $\gamma_G, \gamma_E$                                    | Gaussian kernel coefficients (inverse squared bandwidths) of the two views |
| $Q_{\text{syn}}(t), s_{\text{syn}}(t)$                  | geodesic and Euclidean interpolants of G-SMOTE                             |
| $C_{y_i}$                                               | class-weighted SVM box constraint for class $y_i$                          |
| $F, \bar{m}, \text{SE}$                                 | number of outer folds, fold-mean metric, standard error                    |
| $I_{gk}, \pi_g$                                         | fold-level permutation importance and permutation of family $g$            |
| $h, A, R_{\text{seen}}, R_{\text{new}}, R_{\text{row}}$ | record-level classifier, leakage event, and associated risks               |



**Figure 2.** Schematic illustration of evaluation leakage with repeated recordings. (A) A record-wise split assigns individual recordings to training or test sets, so several subjects (dashed outlines) appear on both sides and the test set is not independent of training identities. (B) A subject-disjoint split assigns all recordings of a subject jointly. (C) Consequence of Equation (18): the record-wise estimate  $R_{\text{row}}$  interpolates between the easier seen-person risk  $R_{\text{seen}}$  and the deployment-relevant unseen-person risk  $R_{\text{new}}$ , drawn for an illustrative  $\Pr(A) = 0.65$ ; the arrow marks the optimistic bias of magnitude  $\Pr(A)(R_{\text{new}} - R_{\text{seen}})$ .

### 2.3. Standardization and Subject Representation


For subject  $i$ , the three acoustic vectors  $x_{ij} \in \mathbb{R}^d$  are standardized using a mean vector  $\mu_k$  and a scale vector estimated exclusively from the recordings of training subjects in outer fold  $k$ . Under the implemented standardization procedure, a zero-variance training coordinate is assigned unit scale:

$$\begin{aligned} \tilde{x}_{ij} &= D_k^{-1}(x_{ij} - \mu_k), D_k = \text{diag}(\sigma_k^*), \\ M_i &= [\tilde{x}_{i1} \tilde{x}_{i2} \tilde{x}_{i3}], \end{aligned} \quad (3)$$

where  $\sigma_k^*$  denotes the training-fold scale vector after zero-variance handling. A thin SVD [57] gives

$$\begin{aligned} M_i &= U_i \Sigma_i V_i^T, Q_i = U_i(:, 1:r), \\ Q_i^T Q_i &= I_r, [Q_i] \in \text{Gr}(r, d). \end{aligned} \quad (4)$$

Throughout, the rank is fixed at  $r=3$ , matching the three repeated recordings per subject. The basis is not unique:  $Q_i$  and  $Q_i O$ , with  $O$  in the orthogonal group  $O(r)$ , represent the same subspace. This quotient structure removes coordinate choices within the retained SVD basis (Figure 3A). In the official releases, one subject per cohort (CONT-36 in UCI-489; subject 37 in PD-252) contains byte-identical repeated recordings, so its matrix has numerical rank two. For such a subject, Equation (4) retains the orthonormal completion supplied by the SVD as the third basis vector, and the reference implementation accepts it only under an explicit opt-in flag.



**Figure 3.** Geometric interpretation of the VoxNeuro representation (schematic; the curved sheet stands for the Grassmann manifold). (A) The three repeated recordings of a subject are mapped by a thin SVD to a single point  $[Q_i]$  on  $\text{Gr}(3, d)$ . (B) The squared chordal distance between two subject subspaces (dashed chord) and the associated Gaussian Grassmann radial-basis-function kernel compare subjects independently of the arbitrary orthonormal basis selected for each subspace; navy and teal points denote subjects of the two classes. (C) A minority

subject is synthesized along the Grassmann geodesic joining two minority neighbors rather than by row-wise Euclidean interpolation.

#### 2.4. Principal Angles and Projection Geometry


Principal angles between two subject subspaces (Figure 4A) are defined by [32,58]

$$Q_i^\top Q_j = A_{ij} \cos(\Theta_{ij}) B_{ij}^\top, \quad \Theta_{ij} = \text{diag}(\theta_{ij,1}, \dots, \theta_{ij,r}), \quad (5)$$

where  $A_{ij}$  and  $B_{ij}$  are orthogonal. With the projector  $P_i = Q_i Q_i^\top$ , the squared chordal distance is given by [33,59]

$$\begin{aligned} d_{\text{ch}}^2([Q_i], [Q_j]) &= \frac{1}{2} \|P_i - P_j\|_F^2 \\ &= r - \|Q_i^\top Q_j\|_F^2 = \sum_{q=1}^r \sin^2 \theta_{ij,q}. \end{aligned} \quad (6)$$

Because  $(Q_i O)(Q_i O)^\top = P_i$ , the distance is basis invariant (Figure 4B). Vectorizing  $P_i$  maps the chordal distance to a Euclidean distance in projector space, up to a constant factor  $\sqrt{2}$  absorbed by the kernel coefficient, thereby providing a Euclidean embedding for the Gaussian Grassmann kernel.



**Figure 4.** The geometry behind Equations (5), (6), (11), and (12), drawn for illustrative toy subspaces. (A) Two planes in  $\mathbb{R}^3$  that share one direction, so the first principal angle is zero and the drawn angle is the second principal angle  $\theta_{ij,2}$ . (B) Exact rank-one relation between a principal angle and the squared chordal distance,  $d_{\text{ch}}^2 = \sin^2 \theta$ . (C) A geodesic path computed with Equations (11) and (12) between two one-dimensional subspaces  $[Q_a]$  and  $[Q_b]$ ; every interpolated frame in the drawn path satisfies  $Q_{\text{syn}}(t)^\top Q_{\text{syn}}(t) = I_r$  to machine precision, consistent with Proposition 2.

#### 2.5. Complementary Euclidean View

The complementary Euclidean view retains the coordinate-wise mean and dispersion:

$$\begin{aligned} \mu_i &= \frac{1}{3} \sum_{j=1}^3 \tilde{x}_{ij}, \sigma_i = \left[ \frac{1}{3} \sum_{j=1}^3 (\tilde{x}_{ij} - \mu_i)^{\odot 2} \right]^{1/2}, \\ s_i &= [\mu_i^\top \sigma_i^\top]^\top \in \mathbb{R}^{2d}, \end{aligned} \quad (7)$$

where  $(\cdot)^{\odot 2}$  and the square root are applied elementwise. Thus  $[Q_i]$  encodes the orientation of the span of the standardized repetitions ( $M_i$  is not centered within subject, so this span generally includes the mean direction), whereas  $s_i$  retains coordinate-level location and dispersion.

#### 2.6. Direct-Sum Kernel and Guarantee of Positive Semidefiniteness

Let  $\phi_G$  and  $\phi_E$  be feature maps for Gaussian kernels on the projection and Euclidean views [38,39]. With fusion weight  $w$ , the joint representation is the Hilbert direct sum

$$\begin{aligned} \Phi(Z_i) &= (\sqrt{w} \phi_G([Q_i]), \sqrt{1-w} \phi_E(s_i)), \\ 0 &\leq w \leq 1, \end{aligned} \quad (8)$$

which induces, for kernel coefficients  $\gamma_G, \gamma_E > 0$  (inverse squared bandwidths),

$$\begin{aligned} K_G(i, j) &= \exp\{-\gamma_G d_{\text{ch}}^2([Q_i], [Q_j])\}, \\ K_E(i, j) &= \exp\{-\gamma_E \|s_i - s_j\|_2^2\}, \\ K(i, j) &= w K_G(i, j) + (1 - w) K_E(i, j). \end{aligned} \quad (9)$$

**Proposition 1** (*basis invariance and positive semidefiniteness*). The kernel in Equation (9) is invariant to  $Q_i \mapsto Q_i O$  and is PSD. Basis invariance follows from the projector representation in Equation (6). Each Gaussian block is PSD on its Euclidean embedding, and a nonnegative weighted sum of PSD kernels is PSD [39,41]; hence, for every  $c \in \mathbb{R}^n$ ,

$$c^T K c = w c^T K_G c + (1 - w) c^T K_E c \geq 0. \quad (10)$$

Equivalently,  $K(i, j) = \langle \Phi(Z_i), \Phi(Z_j) \rangle$  in the direct-sum reproducing-kernel Hilbert space.  $\square$  The projector embedding is essential: Gaussian kernels built on other manifold distances, such as the geodesic arc length, are not PSD for all bandwidths [37], whereas the chordal construction used here inherits positive semidefiniteness from its explicit Euclidean embedding.

### 2.7. Grassmann-Geodesic Synthetic Minority Oversampling

This project-specific manifold adaptation is referred to as Grassmann-geodesic SMOTE (G-SMOTE); the abbreviation does not denote the separate Geometric SMOTE algorithm of Euclidean imbalanced learning [60]. Following the synthetic-neighbor principle of standard SMOTE [44] and its descendants [45,46,61], for minority subjects  $a$  and  $b$ ,

$$\begin{aligned} Q_a^T Q_b &= U \cos(\Theta) V^T, \\ W_{ab} &= (Q_b V - Q_a U \cos \Theta) (\sin \Theta)^\dagger, \end{aligned} \quad (11)$$

where  $\dagger$  denotes the Moore–Penrose pseudoinverse. A shortest geodesic and its coupled Euclidean interpolation are defined as follows (Figure 4C):

$$\begin{aligned} Q_{\text{syn}}(t) &= [Q_a U \cos(t\Theta) + W_{ab} \sin(t\Theta)] V^T, \\ s_{\text{syn}}(t) &= s_a + t(s_b - s_a), \quad t \sim U(0, 1). \end{aligned} \quad (12)$$

Within each imbalanced training fold, the minority class was the control class. For each synthetic subject, one minority subject was sampled uniformly at random, and a partner was selected uniformly from that subject’s five nearest minority neighbors under squared chordal distance. Synthetic subjects were generated until the two class counts were equal. The same sampled value of  $t$  was used for the Grassmann and Euclidean interpolations. For numerical stabilization, the implementation applied a QR factorization to  $Q_{\text{syn}}(t)$ . The resulting invertible right basis change preserves the represented Grassmann point. No synthetic subjects were generated for balanced UCI-489. Interpolating whole-subject representations along the manifold avoids creating synthetic recording rows and selects partners by subspace geometry, but it still interpolates high-dimensional summaries and therefore remains susceptible to the noise-amplification risks of oversampling [47].

**Proposition 2** (*manifold validity*). In exact arithmetic, the nonzero-angle columns of  $Q_a U$  and  $W_{ab}$  are mutually orthonormal, while zero-angle directions remain in the cosine term. Consequently,

$$Q_{\text{syn}}(t)^T Q_{\text{syn}}(t) = V [\cos^2(t\Theta) + \sin^2(t\Theta)] V^T = I_r. \quad (13)$$

Therefore,  $Q_{\text{syn}}(t) \in \text{St}(r, d)$ , the Stiefel manifold of orthonormal  $d \times r$  frames; it represents a valid Grassmann point. Its endpoints are equivalent to  $[Q_a]$  and  $[Q_b]$ .  $\square$

### 2.8. Kernel Support Vector Machine, Metrics, and Complexity

For support vector machine (SVM) training [62], labels are recoded as  $y_i \in \{-1, +1\}$ ,  $Y = \text{diag}(y_1, \dots, y_n)$ , and balanced class weighting sets  $C_{y_i} = Cn / (2n_{y_i})$ , where  $n_{y_i}$  is the number of

217 training subjects in class  $y_i$ ; this is the balanced heuristic implemented in standard open-source libraries  
 218 [63]. When the training classes are exactly balanced, as on UCI-489 and after G-SMOTE parity on  
 219 PD-252, these weights reduce to  $C_{y_i} = C$ ; the weighting is therefore active only for the unaugmented  
 220 class-weighted comparators on imbalanced PD-252. The class-weighted dual and the decision function  
 221 with bias  $b$  are

$$\begin{aligned}
 & \max_{\alpha} \quad 1^T \alpha - \frac{1}{2} \alpha^T Y K Y \alpha, \\
 & \text{subject to} \quad 0 \leq \alpha_i \leq C_{y_i}, y^T \alpha = 0, \\
 & f(Z) = \sum_{i=1}^n \alpha_i y_i K(Z_i, Z) + b.
 \end{aligned} \tag{14}$$

223 Balanced accuracy (BA) [64] and macro-F1 [65] weight diagnostic classes symmetrically:

$$\begin{aligned}
 \text{BA} &= \frac{1}{2} (\text{TPR} + \text{TNR}), \\
 \text{macro-F1} &= \frac{1}{2} (F1_{\text{PD}} + F1_{\text{Ctrl}}),
 \end{aligned} \tag{15}$$

225 where TPR is sensitivity, TNR is specificity, and  $F1_{\text{PD}}$  and  $F1_{\text{Ctrl}}$  are the class-wise F1 scores. The  
 226 Matthews correlation coefficient is a common alternative [66]; balanced accuracy and macro-F1 were  
 227 retained because both weight the two classes symmetrically and are directly interpretable as class-wise  
 228 rates. For  $F=5$  outer folds, a metric  $m$  is summarized by

$$\begin{aligned}
 \bar{m} &= \frac{1}{F} \sum_{k=1}^F m_k, \\
 \text{SE}(\bar{m}) &= \left[ \frac{1}{F(F-1)} \sum_{k=1}^F (m_k - \bar{m})^2 \right]^{1/2}.
 \end{aligned} \tag{16}$$

230 The reported standard error (SE) is a descriptive measure of fold-to-fold variability and is not an  
 231 inferential confidence interval.

232 For  $n$  subjects and rank  $r=3$ , subject SVDs require  $O(n d r^2)$ , the Grassmann Gram matrix requires  
 233  $O(n^2 d r^2)$ , and the Euclidean block requires  $O(n^2 d)$ . These bounds exclude the kernel-SVM  
 234 optimization and the additional G-SMOTE neighbor search, which can dominate as  $n$  grows; all  
 235 components remain tractable for the present cohort sizes  $n=80$  and  $n=252$ .

### 236 2.9. Formal Effect of Record-Wise Leakage

237 The mechanism illustrated in Figure 2 can be stated formally, following record-wise versus subject-wise  
 238 analyses in clinical machine learning [18,20,21]. Let  $(X, y)$  be a random validation recording and its  
 239 label, let  $A$  be the event that the recording belongs to a person represented in training, and let  $h$  be a  
 240 record-level classifier. With  $A^c$  the complement of  $A$ , define

$$\begin{aligned}
 R_{\text{seen}} &= E[\ell(h(X), y) | A], \\
 R_{\text{new}} &= E[\ell(h(X), y) | A^c].
 \end{aligned} \tag{17}$$

242 A record-wise split then estimates

$$\begin{aligned}
 R_{\text{row}} &= \Pr(A) R_{\text{seen}} + [1 - \Pr(A)] R_{\text{new}}, \\
 R_{\text{row}} - R_{\text{new}} &= \Pr(A) (R_{\text{seen}} - R_{\text{new}}).
 \end{aligned} \tag{18}$$

244 If repeated recordings from previously observed people are easier to classify than recordings from  
 245 unseen people, then  $R_{\text{seen}} < R_{\text{new}}$ , and the record-wise estimate is optimistically biased in direct  
 246 proportion to  $\Pr(A)$  (Figure 2C). Subject-disjoint folds enforce  $\Pr(A)=0$  by construction.




## 2.10. Experimental Protocol

Outer folds were generated by five-fold stratified subject-level cross-validation with shuffling and a random seed of 42 (Figure 5A). For each outer fold, the recording-level standardizer was fitted only to recordings from training subjects. Each kernel coefficient was estimated from the positive squared pairwise distances  $\delta^2$  of the corresponding view over the training subjects, after G-SMOTE augmentation on PD-252, as  $\gamma = [2 \text{ median}(\delta^2) + 10^{-12}]^{-1}$ . Rank  $r = 3$ , fusion weight  $w = 0.5$ , and SVM penalty  $C = 1$  were fixed across datasets; no per-dataset hyperparameter search was performed, thereby avoiding reuse of the outer test folds for hyperparameter selection [23,24]. Because UCI-489 is balanced, no oversampling was applied there. On PD-252, the comparator set additionally included SMOTE plus logistic regression and a matched Euclidean RBF SVM using identical augmentation restricted to the training fold, so that the fused model and its nonlinear Euclidean ablation differed only in the Grassmann kernel term (Figure 5B). All methods were implemented in Python, and the implementation is available in the project repository (see the Data Availability Statement). Table 3 summarizes all fixed protocol choices. All comparators were implemented in the released codebase and verified against their stated configurations; the term audited comparators refers to this set, as distinct from values reported by other studies.

The comparator set was designed so that each comparison isolates one design decision. Class-weighted logistic regression and class-weighted linear SVM probe whether simple decision boundaries on the Euclidean summary already suffice. SMOTE plus logistic regression tests a standard imbalanced-learning remedy [44]. The matched Euclidean RBF SVM with identical G-SMOTE augmentation is the critical ablation: it shares the summary features, the resampling, and the kernel family of the fused model while omitting only the Grassmann kernel term; the shared augmentation itself selects partners by chordal distance (Section 2.7). Any advantage of the fused model over this ablation is therefore attributable to the geometric kernel component, conditional on the shared augmentation, rather than to nonlinearity or resampling. On PD-252, G-SMOTE randomly samples anchor subjects and neighbors and draws interpolation coefficients, so performance varies with the augmentation seed. To isolate this source of variability, the fixed five-fold evaluation was repeated for 24 distinct augmentation seeds while the outer-fold assignments, preprocessing, hyperparameters, and all non-augmentation settings were held fixed. Within each seed and fold, the fused model and the matched Euclidean RBF ablation used the same augmented training set, so their difference is paired at the seed level. The random state of the original PD-252 augmentation realization was not retained, so that specific run cannot be regenerated exactly from a seed; the seed analysis of Section 3.3 therefore characterizes where it lies within the distribution of realizations.

**Table 3.** Fixed experimental protocol. All learned components are fitted inside the outer training folds only.

| Component               | Setting                                                                                                                                      |
|-------------------------|----------------------------------------------------------------------------------------------------------------------------------------------|
| Outer evaluation        | five-fold stratified subject-level cross-validation, shuffled, seed 42                                                                       |
| Standardization         | per-coordinate mean/scale from training recordings; zero-variance coordinates assigned unit scale                                            |
| Subject representation  | rank $r = 3$ thin SVD (Grassmann view) + mean/dispersion vector (Euclidean view)                                                             |
| Kernel coefficients     | $\gamma = [2 \text{ median}(\delta^2) + 10^{-12}]^{-1}$ from positive squared pairwise training distances (augmented training set on PD-252) |
| Fusion weight / penalty | $w = 0.5$ ; class-weighted SVM with $C = 1$ ; no per-dataset tuning                                                                          |
| Oversampling            | none on UCI-489; G-SMOTE to class parity on PD-252, five nearest minority neighbors under chordal distance                                   |
| Comparators             | class-weighted logistic regression; class-weighted linear SVM; SMOTE + logistic regression and matched Euclidean RBF SVM on PD-252           |
| Metrics                 | balanced accuracy and macro-F1, mean $\pm$ fold-level SE; pooled out-of-fold confusion counts                                                |



**Figure 5.** Evaluation protocol. (A) Five-fold subject-disjoint partition; in each outer fold one subject block (orange) is held out and never influences any fitted component. (B) Everything that learns from data—the standardizer, G-SMOTE, the kernel coefficients, and the SVM—is fitted inside the training fold only. (C) Exactly one prediction is issued per held-out subject, fold metrics are computed at the subject level, and Equation (16) aggregates them descriptively.

### 2.11. Feature-Family Permutation Sensitivity

For feature family  $g$ , the fold-level decrease in performance and its summary are

$$I_{gk} = m_k - m_k^{\pi_g}, \hat{I}_g = \frac{1}{F} \sum_{k=1}^F I_{gk},$$


$$\text{SE}(\hat{I}_g) = \left[ \frac{1}{F(F-1)} \sum_{k=1}^F (I_{gk} - \hat{I}_g)^2 \right]^{1/2}.$$

Sensitivity of the Euclidean summary was evaluated on each held-out fold without refitting. The coordinates of  $s_i$  associated with family  $g$  were jointly permuted across held-out subjects while  $Q_i$  and all other summary coordinates were held fixed. In Equation (19),  $m_k$  is the original held-out macro-F1 and  $m_k^{\pi_g}$  is macro-F1 after permutation. Each family and fold used a single permutation realization, so fold-level variability also includes permutation-draw noise.

## 3. Results

### 3.1. Performance Across Datasets

Figure 6 shows that the fused model attained the highest mean balanced accuracy and macro-F1 among the audited comparators in both datasets; Table 4 summarizes the principal scores. On UCI-489, Grassmann–Euclidean fusion achieved a balanced accuracy of 0.850 with an SE of 0.042 and a macro-F1 of 0.847 with an SE of 0.043. Its mean scores were 0.075 balanced-accuracy points and 0.077 macro-F1 points higher than those of class-weighted logistic regression, the strongest UCI-489 baseline. On PD-252, fusion with G-SMOTE achieved a balanced accuracy of 0.749 with an SE of 0.039 and a macro-F1 of 0.760 with an SE of 0.035. The matched Euclidean RBF model with identical augmentation achieved a balanced accuracy of 0.736 with an SE of 0.035 and a macro-F1 of 0.735 with an SE of 0.027. Thus, the fused model’s mean scores were 0.013 balanced-accuracy points and 0.025 macro-F1 points higher than those of the matched nonlinear Euclidean ablation. These are descriptive cross-validation differences from a single G-SMOTE realization on PD-252 (Section 3.3); no inferential test of model superiority was performed.




**Figure 6.** Subject-level speech results for (A) class-balanced UCI-489 and (B) high-dimensional, class-imbalanced PD-252. Whiskers denote standard error across five outer folds. RBF + G-SMOTE is the matched nonlinear Euclidean ablation and SMOTE + LogReg is SMOTE with logistic regression; Balanced LogReg and Logistic denote class-weighted logistic regression, and Balanced linear SVM and Linear SVM denote class-weighted linear SVM. PD-252 values for the fused model and the matched Euclidean ablation are from the original single G-SMOTE augmentation realization; paired 24-seed summaries are reported in Table 6.

**Table 4.** Headline subject-level performance (mean  $\pm$  fold-level standard error over five subject-disjoint folds). The strongest audited baseline is shown for each dataset; all comparator summary estimates appear in Figure 6. PD-252 rows report the original single G-SMOTE realization; seed sensitivity is quantified in Table 6.

| Dataset | Model                                      | Balanced acc.     | Macro-F1          |
|---------|--------------------------------------------|-------------------|-------------------|
| UCI-489 | Grassmann-Euclidean fusion                 | $0.850 \pm 0.042$ | $0.847 \pm 0.043$ |
| UCI-489 | Class-weighted logistic regression         | $0.775 \pm 0.042$ | $0.770 \pm 0.043$ |
| PD-252  | Fusion + G-SMOTE                           | $0.749 \pm 0.039$ | $0.760 \pm 0.035$ |
| PD-252  | Euclidean RBF + G-SMOTE (matched ablation) | $0.736 \pm 0.035$ | $0.735 \pm 0.027$ |

### 3.2. Operating Characteristics

The UCI-489 fused model produced 34 true-negative (control), six false-positive, six false-negative, and 34 true-positive (PD) classifications, yielding sensitivity and specificity of 0.850. On PD-252, under the original single G-SMOTE realization, the pooled confusion matrix contained 173 true-positive (PD), 15 false-negative, 37 true-negative (control), and 27 false-positive classifications, corresponding to sensitivity 0.920, specificity 0.578, positive predictive value 0.865, and negative predictive value 0.712 (Figure 7, Table 5). Sensitivity and specificity are prevalence-free, whereas PPV and NPV depend on this cohort’s 188/64 composition and the default decision threshold and do not transport to lower-prevalence screening settings. Rates computed from pooled confusion matrices are descriptive and need not equal the arithmetic mean of fold-level metrics. The PD-252 profile is sensitive but would create a substantial referral burden. This operating profile warrants further evaluation of the method as a research-stage screening aid and does not support its use for autonomous diagnosis.



**Figure 7.** Pooled out-of-fold confusion matrices for (A) UCI-489 and (B) PD-252, containing one prediction per subject; percentages are normalized within the true class. PD-252 counts are from the original single G-SMOTE augmentation realization.

**Table 5.** Pooled out-of-fold operating characteristics computed from the confusion counts of Figure 7 (UCI-489: TP = 34, FN = 6, TN = 34, FP = 6; PD-252: TP = 173, FN = 15, TN = 37, FP = 27). TP, true positive; FN, false negative; TN, true negative; FP, false positive; PPV, positive predictive value; NPV, negative predictive value.

| Dataset | Sensitivity | Specificity | PPV   | NPV   |
|---------|-------------|-------------|-------|-------|
| UCI-489 | 0.850       | 0.850       | 0.850 | 0.850 |
| PD-252  | 0.920       | 0.578       | 0.865 | 0.712 |

### 3.3. Sensitivity to the G-SMOTE Augmentation Seed

The PD-252 scores in Table 4 were obtained from one G-SMOTE augmentation realization. The paired analysis of Section 2.10 characterizes the sensitivity of both models and of their difference to the augmentation seed (Table 6). Across 24 seed runs, balanced accuracy was  $0.739 \pm 0.012$  for the fused model and  $0.731 \pm 0.013$  for the matched Euclidean ablation, and macro-F1 was  $0.752 \pm 0.013$  and  $0.729 \pm 0.013$ , respectively (mean  $\pm$  standard deviation across seeds). The paired macro-F1 difference averaged  $+0.023 \pm 0.015$  and was positive in 23 of 24 runs; the original realization’s difference of  $+0.025$  lies within this distribution. The paired balanced-accuracy difference averaged  $+0.009 \pm 0.014$ , was positive in 16 of 24 runs, and reversed sign in 8 of 24. The original model scores lie within their corresponding empirical seed ranges, 0.6–0.8 seed standard deviations above the seed means, and the reference implementation’s default seed yields 0.731 and 0.749 for the fused model, likewise within these ranges; this run and its out-of-fold predictions are archived in the project repository. Under seed means, the fused model again attained the highest mean balanced accuracy and macro-F1 among the audited PD-252 comparators; the two class-weighted comparators without augmentation are unaffected by this seed, and SMOTE plus logistic regression is represented by its single archived realization. Thus, conditional on the fixed outer folds, the macro-F1 margin was directionally consistent in all but one seed run, whereas the balanced-accuracy margin can be reversed by individual draws.


**Table 6.** Sensitivity of the PD-252 comparison to the G-SMOTE augmentation seed: paired evaluation of the fused model and the matched Euclidean RBF ablation across 24 distinct seeds with outer folds, preprocessing, and hyperparameters held fixed. Values are means  $\pm$  standard deviations (SD) across seeds;  $\Delta$  denotes the paired difference (fused minus ablation) computed within each seed, summarized by its mean  $\pm$  SD, median [interquartile range, IQR], and sign count.

| Metric            | Fused             | Ablation          | $\Delta$ mean $\pm$ SD | $\Delta$ median [IQR]     | $\Delta > 0$ |
|-------------------|-------------------|-------------------|------------------------|---------------------------|--------------|
| Balanced accuracy | $0.739 \pm 0.012$ | $0.731 \pm 0.013$ | $+0.009 \pm 0.014$     | $+0.008 [-0.002, +0.018]$ | 16/24        |
| Macro-F1          | $0.752 \pm 0.013$ | $0.729 \pm 0.013$ | $+0.023 \pm 0.015$     | $+0.020 [+0.012, +0.034]$ | 23/24        |

Each seed-level value is the five-fold cross-validation point estimate under fixed outer folds; standard deviations describe variation across the 24 augmentation seeds and are distinct from the fold-level standard errors of Table 4.  $\Delta$  ranges: balanced accuracy  $[-0.013, +0.042]$ ; macro-F1  $[-0.001, +0.061]$ . Per-seed values are archived in the project repository.

### 3.4. Feature-Family Sensitivity

The TQWT summary block produced the largest mean decrease in held-out macro-F1, at 0.143 with an SE of 0.037, followed by the MFCC summary block, at 0.050 with an SE of 0.017 (Figure 8). The vocal/time-frequency, wavelet, IMF/EMD, and demographic families produced mean decreases near zero. These exploratory values measure dependence of the fitted model on the Euclidean summary coordinates under correlated features; they do not establish biological causality and, because only the Euclidean view is permuted, do not quantify the dependence of the full two-view model on each family. The prominence of the TQWT family is consistent with the descriptor engineering of the source dataset, where tunable Q-factor features were introduced precisely for Parkinsonian voice [31,55].



**Figure 8.** PD-252 feature-family permutation sensitivity on the Euclidean summary view. Positive values denote a decrease in held-out macro-F1 after permutation, and whiskers denote fold-level standard error. TQWT denotes tunable Q-factor wavelet transform; MFCC denotes Mel-frequency cepstral coefficients; IMF/EMD denotes intrinsic mode function/empirical mode decomposition; Demographics denotes the dataset-provided sex/gender covariate.

## 4. Discussion


The central methodological contribution is the alignment of the evaluation unit with the clinical unit. Subject-disjoint validation asks whether a model generalizes to an unseen person. The Grassmann representation uses the dependence among repeated recordings rather than treating the recordings as independent replicates. Basis invariance removes arbitrary coordinate choices, while the Euclidean view restores the coordinate-level location and dispersion magnitudes that a pure subspace representation discards.

Across both datasets, the fused model attained the highest mean balanced accuracy and macro-F1 among the audited comparators, and on PD-252 it exceeded the matched Euclidean RBF ablation with G-SMOTE. This establishes the highest mean performance within the audited comparator set; it does not establish universal or statistically significant superiority over all learning algorithms. The seed-sensitivity analysis of Section 3.3 sharpens the ablation comparison: the paired macro-F1 difference was positive in 23 of 24 augmentation runs, whereas the balanced-accuracy difference was positive in 16 of 24 and reversed in 8. These counts characterize sensitivity to the augmentation randomization under fixed outer folds; they are not population-level inference about model superiority. With respect to the


two objectives of Section 1, the second is answered by this paired comparison: the macro-F1 advantage was consistent across seeds, whereas the balanced-accuracy advantage was less stable. The first was addressed analytically, without a matched record-wise experiment: Equations (17) and (18) make the mechanism behind the design explicit. Whenever repeated recordings from previously observed people are easier to classify than recordings from unseen people, record-wise scores are optimistically biased in direct proportion to the probability that a validation recording belongs to an already-observed person. Subject-disjoint folds remove that term entirely. Reported subject-level scores are therefore expected to be lower than record-wise figures reported on the same corpora, but they answer the question that a screening deployment poses [20,22].

An evaluation-focused perspective also clarifies how these results should be interpreted in relation to the wider literature. Reported accuracies for voice-based PD detection vary widely across studies, and validation strategy is a recognized source of that variation [16,17]. Record-wise splitting remains common in secondary analyses of repeated-recording corpora, whereas replication-aware methods have also been developed for UCI-489 [29]. Balanced accuracy and macro-F1 were used because both cohorts were evaluated at the subject level and PD-252 had unequal class sizes. Both metrics weight the diagnostic classes symmetrically rather than rewarding majority-class predictions [64,65]. Fold-level standard errors are reported descriptively, and no significance test is claimed: five outer folds provide a descriptive summary of variability, not a basis for inferential comparison [25].

The deployed VoxNeuro web application (Figure 9) guides three speech tasks (sustained vowel, read passage, and free speech), records them with only a commodity microphone and a web browser, and visualizes session scores over time, which keeps the barrier to repeated assessment low. Figure 10 places the prototype in an envisioned deployment workflow and distinguishes the components evaluated here from capture and decision-support stages that were not evaluated. This is precisely the kind of low-friction, high-frequency measurement that technology assessments of PD care have called for [50,51]. It is a research-stage prototype, not an autonomous diagnostic system or medical device, and no application-collected data were used in the analyses reported here. Prospective use requires explicit privacy, consent, security, accessibility conformance, calibration, and clinical governance procedures.



**Figure 9.** Deployed VoxNeuro web application (research-stage prototype). (A) Home view listing the three guided speech tasks. (B) Sustained-vowel capture screen. (C) Read-passages capture screen, which uses the standard Rainbow Passage. (D) Longitudinal progress view with illustrative entries. The application requires only a commodity microphone and a web browser; no application-collected data were used in the analyses reported here.



**Figure 10.** Envisioned deployment workflow of the VoxNeuro prototype. Solid boxes mark the components evaluated in this study on the two public cohorts (subject representation, kernel fusion with G-SMOTE, and subject-level scoring); dashed boxes mark capture and decision-support stages of the research-stage prototype that were not evaluated here. Descriptor extraction in this study used the dataset-provided engineered features, and any referral decision remains a clinical judgement.



This study has limitations. The speech cohorts contain retrospective engineered features rather than prospectively acquired raw audio. The study included no probability calibration [67,68], decision-curve analysis [69], independent clinical test cohort, device-shift evaluation, language-shift audit, or prespecified subgroup analysis. The fusion weight and SVM penalty were fixed rather than tuned, so sensitivity to  $w$  was not characterized, and the comparator set contained no deep or pretrained audio models, which the engineered-feature cohorts cannot support. Because high PD-252 sensitivity coexists with modest specificity, a credible system should support referral and longitudinal review rather than replace neurological assessment. Future work should preregister preprocessing and thresholds and report in accordance with structured guidance for clinical prediction models [70,71]; collect repeated raw speech prospectively; and evaluate calibration, test–retest reliability, fairness, and clinician-facing decision consequences. The two-view construction extends naturally to additional recording tasks and views through further positive-semidefinite kernel terms, with weights selected by nested multiple-kernel optimization [42,43].

Reproducibility safeguards included fixed subject-disjoint folds, training-fold-only preprocessing, fixed hyperparameters, saved fold-level metrics, pooled confusion counts, and restriction of reported results to audited comparator implementations. The project repository archives subject-level out-of-fold predictions with fold identifiers for UCI-489, a pinned dependency environment, and the seed-sensitivity script with its 24 per-seed paired metrics. The feature-family permutation analysis of Section 3.4 is reported from the original study artifacts and is not regenerable from the repository. The UCI-489 aggregates of Tables 4 and 5 can therefore be regenerated exactly and the seed-level summaries of Table 6 reconstructed; the original PD-252 realization is situated within the archived seed distribution (Section 3.3).

## 5. Conclusions

VoxNeuro provides a leakage-safe framework for repeated Parkinsonian speech. Each subject is represented by a Grassmann subspace and a Euclidean mean–variability vector. The corresponding positive-semidefinite kernels are fused, and minority subjects are synthesized along manifold geodesics only within training folds. The fused model attained the highest mean balanced accuracy and macro-F1 among the audited comparators on UCI-489 (0.850 and 0.847) and, across 24 paired augmentation seeds, on PD-252 ( $0.739 \pm 0.012$  and  $0.752 \pm 0.013$ ). Its macro-F1 exceeded that of the matched Euclidean RBF ablation in 23 of 24 draws. The deployed web application implements guided capture and longitudinal visualization as a research-stage prototype. Clinically credible digital biomarkers require auditable implementations, transparent uncertainty, and evaluation designs that test generalization to unseen people; prospective external validation is the necessary next step before any clinical use.

**Author Contributions:** Conceptualization, L.K.; methodology, L.K.; software, L.K.; validation, L.K. and M.W.; formal analysis, L.K.; investigation, L.K. and M.W.; data curation, L.K.; writing—original draft preparation, L.K.; writing—review and editing, L.K. and M.W.; visualization, L.K.; project administration, L.K. All authors have read and agreed to the published version of the manuscript.

**Funding:** This research received no external funding.

**Institutional Review Board Statement:** Ethical review and approval were not required for this study, as it consisted exclusively of a secondary computational analysis of publicly available, fully de-identified datasets; no participants were recruited, and no new human data were collected.

**Informed Consent Statement:** Not applicable. The study analyzed publicly available, de-identified datasets and did not involve direct human participation; informed consent and ethics procedures for the original data collections are documented in the source publications cited in Section 2.1.

**Data Availability Statement:** The datasets analyzed in this study are openly available from the UCI Machine Learning Repository: the Parkinson Dataset with Replicated Acoustic Features (UCI-489; UCI dataset 489, doi:10.24432/C5701F) and the Parkinson’s Disease Classification dataset (PD-252; UCI dataset 470, doi:10.24432/C5MS4X). Source code is available at <https://github.com/khan-laiba/voxneuro> (release v1.0.0); the repository also archives subject-level out-of-fold prediction files for UCI-489, the pinned dependency environment,



and the per-seed outputs of the PD-252 seed-sensitivity analysis. The deployed research-stage web application is available at <https://voxneuro.org>. No new primary data were collected in this study.

**Acknowledgments:** During the preparation of this manuscript, the authors used Claude (Anthropic; Claude Opus) to draft and edit manuscript text, format the manuscript, and produce the schematic illustrations in Figures 1, 2, 3, 4, 5, 7, and 10, and assemble the screenshot panels of Figure 9, and GPT-5.6 Sol (OpenAI; via the Codex CLI) to review manuscript language and figure design. The authors have reviewed and edited the output and take full responsibility for the content of this publication.

**Conflicts of Interest:** The authors declare no conflicts of interest.

## References

1. Bloem, B.R.; Okun, M.S.; Klein, C. Parkinson's disease. *Lancet* **2021**, *397*, 2284–2303.
2. Kalia, L.V.; Lang, A.E. Parkinson's disease. *Lancet* **2015**, *386*, 896–912.
3. Poewe, W.; Seppi, K.; Tanner, C.M.; Halliday, G.M.; Brundin, P.; Volkmann, J.; Schrag, A.-E.; Lang, A.E. Parkinson disease. *Nat. Rev. Dis. Primers* **2017**, *3*, 17013.
4. Postuma, R.B.; Berg, D.; Stern, M.; Poewe, W.; Olanow, C.W.; Oertel, W.; Obeso, J.; Marek, K.; Litvan, I.; Lang, A.E.; et al. MDS clinical diagnostic criteria for Parkinson's disease. *Mov. Disord.* **2015**, *30*, 1591–1601.
5. Armstrong, M.J.; Okun, M.S. Diagnosis and treatment of Parkinson disease: A review. *JAMA* **2020**, *323*, 548–560.
6. GBD 2016 Parkinson's Disease Collaborators. Global, regional, and national burden of Parkinson's disease, 1990–2016: A systematic analysis for the Global Burden of Disease Study 2016. *Lancet Neurol.* **2018**, *17*, 939–953.
7. Dorsey, E.R.; Bloem, B.R. The Parkinson pandemic—A call to action. *JAMA Neurol.* **2018**, *75*, 9–10.
8. Logemann, J.A.; Fisher, H.B.; Boshes, B.; Blonsky, E.R. Frequency and cooccurrence of vocal tract dysfunctions in the speech of a large sample of Parkinson patients. *J. Speech Hear. Disord.* **1978**, *43*, 47–57.
9. Ho, A.K.; Iannsek, R.; Marigliani, C.; Bradshaw, J.L.; Gates, S. Speech impairment in a large sample of patients with Parkinson's disease. *Behav. Neurol.* **1998**, *11*, 131–137.
10. Rusz, J.; Cmejla, R.; Ruzickova, H.; Ruzicka, E. Quantitative acoustic measurements for characterization of speech and voice disorders in early untreated Parkinson's disease. *J. Acoust. Soc. Am.* **2011**, *129*, 350–367.
11. Little, M.A.; McSharry, P.E.; Roberts, S.J.; Costello, D.A.E.; Moroz, I.M. Exploiting nonlinear recurrence and fractal scaling properties for voice disorder detection. *Biomed. Eng. Online* **2007**, *6*, 23.
12. Little, M.A.; McSharry, P.E.; Hunter, E.J.; Spielman, J.; Ramig, L.O. Suitability of dysphonia measurements for telemonitoring of Parkinson's disease. *IEEE Trans. Biomed. Eng.* **2009**, *56*, 1015–1022.
13. Tsanas, A.; Little, M.A.; McSharry, P.E.; Ramig, L.O. Accurate telemonitoring of Parkinson's disease progression by noninvasive speech tests. *IEEE Trans. Biomed. Eng.* **2010**, *57*, 884–893.
14. Tsanas, A.; Little, M.A.; McSharry, P.E.; Spielman, J.; Ramig, L.O. Novel speech signal processing algorithms for high-accuracy classification of Parkinson's disease. *IEEE Trans. Biomed. Eng.* **2012**, *59*, 1264–1271.
15. Hlavnička, J.; Čmejla, R.; Tykalová, T.; Šonka, K.; Růžička, E.; Rusz, J. Automated analysis of connected speech reveals early biomarkers of Parkinson's disease in patients with rapid eye movement sleep behaviour disorder. *Sci. Rep.* **2017**, *7*, 12.
16. Sedigh Malekroodi, H.; Lee, B.-i.; Yi, M. Voice-based detection of Parkinson's disease using machine and deep learning approaches: A systematic review. *Bioengineering* **2025**, *12*, 1279.
17. Moro-Velazquez, L.; Gomez-Garcia, J.A.; Arias-Londoño, J.D.; Dehak, N.; Godino-Llorente, J.I. Advances in Parkinson's disease detection and assessment using voice and speech: A review of the state-of-the-art. *Biomed. Signal Process. Control* **2021**, *66*, 102418.
18. Kaufman, S.; Rosset, S.; Perlich, C.; Stitelman, O. Leakage in data mining: Formulation, detection, and avoidance. *ACM Trans. Knowl. Discov. Data* **2012**, *6*, 15.
19. Kapoor, S.; Narayanan, A. Leakage and the reproducibility crisis in machine-learning-based science. *Patterns* **2023**, *4*, 100804.
20. Saeb, S.; Lonini, L.; Jayaraman, A.; Mohr, D.C.; Kording, K.P. The need to approximate the use-case in clinical machine learning. *Gigascience* **2017**, *6*, gix019.
21. Little, M.A.; Varoquaux, G.; Saeb, S.; Lonini, L.; Jayaraman, A.; Mohr, D.C.; Kording, K.P. Using and understanding cross-validation strategies. Perspectives on Saeb et al. *Gigascience* **2017**, *6*, gix020.
22. Chaibub Neto, E.; Pratap, A.; Perumal, T.M.; Tummacherla, M.; Snyder, P.; Bot, B.M.; Trister, A.D.; Friend, S.H.; Mangravite, L.; Omberg, L. Detecting the impact of subject characteristics on machine learning-based diagnostic applications. *NPJ Digit. Med.* **2019**, *2*, 99.
23. Cawley, G.C.; Talbot, N.L.C. On over-fitting in model selection and subsequent selection bias in performance evaluation. *J. Mach. Learn. Res.* **2010**, *11*, 2079–2107.

537 24. Varma, S.; Simon, R. Bias in error estimation when using cross-validation for model selection. *BMC*  
538 *Bioinformatics* **2006**, *7*, 91.

539 25. Varoquaux, G. Cross-validation failure: Small sample sizes lead to large error bars. *NeuroImage* **2018**, *180*, 68–77.

540 26. Arlot, S.; Celisse, A. A survey of cross-validation procedures for model selection. *Stat. Surv.* **2010**, *4*, 40–79.

541 27. Hastie, T.; Tibshirani, R.; Friedman, J. *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*,  
542 2nd ed.; Springer: New York, NY, USA, 2009.

543 28. Pérez, C.J. *Parkinson Dataset with Replicated Acoustic Features*; UCI Machine Learning Repository, 2016.  
544 doi:10.24432/C5701F.

545 29. Naranjo, L.; Pérez, C.J.; Martín, J.; Campos-Roca, Y. A two-stage variable selection and classification approach  
546 for Parkinson's disease detection by using voice recording replications. *Comput. Methods Programs Biomed.* **2017**,  
547 *142*, 147–156.

548 30. Sakar, B.E.; Isenkol, M.E.; Sakar, C.O.; Sertbas, A.; Gurgun, F.; Delil, S.; Apaydin, H.; Kursun, O. Collection and  
549 analysis of a Parkinson speech dataset with multiple types of sound recordings. *IEEE J. Biomed. Health Inform.*  
550 **2013**, *17*, 828–834.

551 31. Sakar, C.O.; Serbes, G.; Gunduz, A.; Tunc, H.C.; Nizam, H.; Sakar, B.E.; Tutuncu, M.; Aydin, T.; Isenkol, M.E.;  
552 Apaydin, H. A comparative analysis of speech signal processing algorithms for Parkinson's disease  
553 classification and the use of the tunable Q-factor wavelet transform. *Appl. Soft Comput.* **2019**, *74*, 255–263.

554 32. Wong, Y.-C. Differential geometry of Grassmann manifolds. *Proc. Natl. Acad. Sci. USA* **1967**, *57*, 589–594.

555 33. Edelman, A.; Arias, T.A.; Smith, S.T. The geometry of algorithms with orthogonality constraints. *SIAM J. Matrix*  
556 *Anal. Appl.* **1998**, *20*, 303–353.

557 34. Absil, P.-A.; Mahony, R.; Sepulchre, R. *Optimization Algorithms on Matrix Manifolds*; Princeton University Press:  
558 Princeton, NJ, USA, 2008.

559 35. Hamm, J.; Lee, D.D. Grassmann discriminant analysis: A unifying view on subspace-based learning. In  
560 Proceedings of the 25th International Conference on Machine Learning, Helsinki, Finland, 5–9 July 2008; pp.  
561 376–383.

562 36. Turaga, P.; Veeraraghavan, A.; Srivastava, A.; Chellappa, R. Statistical computations on Grassmann and Stiefel  
563 manifolds for image and video-based recognition. *IEEE Trans. Pattern Anal. Mach. Intell.* **2011**, *33*, 2273–2286.

564 37. Jayasumana, S.; Hartley, R.; Salzmann, M.; Li, H.; Harandi, M. Kernel methods on Riemannian manifolds with  
565 Gaussian RBF kernels. *IEEE Trans. Pattern Anal. Mach. Intell.* **2015**, *37*, 2464–2477.

566 38. Aronszajn, N. Theory of reproducing kernels. *Trans. Am. Math. Soc.* **1950**, *68*, 337–404.

567 39. Schölkopf, B.; Smola, A.J. *Learning with Kernels: Support Vector Machines, Regularization, Optimization, and Beyond*;  
568 MIT Press: Cambridge, MA, USA, 2002.

569 40. Shawe-Taylor, J.; Cristianini, N. *Kernel Methods for Pattern Analysis*; Cambridge University Press: Cambridge,  
570 UK, 2004.

571 41. Lanckriet, G.R.G.; Cristianini, N.; Bartlett, P.; El Ghaoui, L.; Jordan, M.I. Learning the kernel matrix with  
572 semidefinite programming. *J. Mach. Learn. Res.* **2004**, *5*, 27–72.

573 42. Gönen, M.; Alpaydin, E. Multiple kernel learning algorithms. *J. Mach. Learn. Res.* **2011**, *12*, 2211–2268.

574 43. Rakotomamonjy, A.; Bach, F.R.; Canu, S.; Grandvalet, Y. SimpleMKL. *J. Mach. Learn. Res.* **2008**, *9*, 2491–2521.

575 44. Chawla, N.V.; Bowyer, K.W.; Hall, L.O.; Kegelmeyer, W.P. SMOTE: Synthetic minority over-sampling  
576 technique. *J. Artif. Intell. Res.* **2002**, *16*, 321–357.

577 45. He, H.; Garcia, E.A. Learning from imbalanced data. *IEEE Trans. Knowl. Data Eng.* **2009**, *21*, 1263–1284.

578 46. Fernández, A.; García, S.; Herrera, F.; Chawla, N.V. SMOTE for learning from imbalanced data: Progress and  
579 challenges, marking the 15-year anniversary. *J. Artif. Intell. Res.* **2018**, *61*, 863–905.

580 47. Blagus, R.; Lusa, L. SMOTE for high-dimensional class-imbalanced data. *BMC Bioinformatics* **2013**, *14*, 106.

581 48. Bot, B.M.; Suver, C.; Neto, E.C.; Kellen, M.; Klein, A.; Bare, C.; Doerr, M.; Pratap, A.; Wilbanks, J.; Dorsey, E.R.;  
582 et al. The mPower study, Parkinson disease mobile data collected using ResearchKit. *Sci. Data* **2016**, *3*, 160011.

583 49. Arora, S.; Venkataraman, V.; Zhan, A.; Donohue, S.; Biglan, K.M.; Dorsey, E.R.; Little, M.A. Detecting and  
584 monitoring the symptoms of Parkinson's disease using smartphones: A pilot study. *Parkinsonism Relat. Disord.*  
585 **2015**, *21*, 650–653.

586 50. Espay, A.J.; Bonato, P.; Nahab, F.B.; Maetzler, W.; Dean, J.M.; Klucken, J.; Eskofier, B.M.; Merola, A.; Horak, F.;  
587 Lang, A.E.; et al. Technology in Parkinson's disease: Challenges and opportunities. *Mov. Disord.* **2016**, *31*, 1272–  
588 1282.

589 51. Dorsey, E.R.; Glidden, A.M.; Holloway, M.R.; Birbeck, G.L.; Schwamm, L.H. Teleneurology and mobile  
590 technologies: The future of neurological care. *Nat. Rev. Neurol.* **2018**, *14*, 285–297.

591 52. Topol, E.J. High-performance medicine: The convergence of human and artificial intelligence. *Nat. Med.* **2019**, *25*,  
592 44–56.

53. Sakar, C.O.; Serbes, G.; Gunduz, A.; Nizam, H.; Sakar, B.E. *Parkinson's Disease Classification*; UCI Machine Learning Repository, 2018. doi:10.24432/C5MS4X.
54. Davis, S.B.; Mermelstein, P. Comparison of parametric representations for monosyllabic word recognition in continuously spoken sentences. *IEEE Trans. Acoust. Speech Signal Process.* **1980**, *28*, 357–366.
55. Selesnick, I.W. Wavelet transform with tunable Q-factor. *IEEE Trans. Signal Process.* **2011**, *59*, 3560–3575.
56. Huang, N.E.; Shen, Z.; Long, S.R.; Wu, M.C.; Shih, H.H.; Zheng, Q.; Yen, N.-C.; Tung, C.C.; Liu, H.H. The empirical mode decomposition and the Hilbert spectrum for nonlinear and non-stationary time series analysis. *Proc. R. Soc. Lond. A* **1998**, *454*, 903–995.
57. Golub, G.H.; Van Loan, C.F. *Matrix Computations*, 4th ed.; Johns Hopkins University Press: Baltimore, MD, USA, 2013.
58. Björck, Å.; Golub, G.H. Numerical methods for computing angles between linear subspaces. *Math. Comput.* **1973**, *27*, 579–594.
59. Chikuse, Y. *Statistics on Special Manifolds*; Springer: New York, NY, USA, 2003.
60. Douzas, G.; Bação, F. Geometric SMOTE: A geometrically enhanced drop-in replacement for SMOTE. *Inf. Sci.* **2019**, *501*, 118–135.
61. Batista, G.E.A.P.A.; Prati, R.C.; Monard, M.C. A study of the behavior of several methods for balancing machine learning training data. *ACM SIGKDD Explor. Newsl.* **2004**, *6*, 20–29.
62. Cortes, C.; Vapnik, V. Support-vector networks. *Mach. Learn.* **1995**, *20*, 273–297.
63. Pedregosa, F.; Varoquaux, G.; Gramfort, A.; Michel, V.; Thirion, B.; Grisel, O.; Blondel, M.; Prettenhofer, P.; Weiss, R.; Dubourg, V.; et al. Scikit-learn: Machine learning in Python. *J. Mach. Learn. Res.* **2011**, *12*, 2825–2830.
64. Brodersen, K.H.; Ong, C.S.; Stephan, K.E.; Buhmann, J.M. The balanced accuracy and its posterior distribution. In Proceedings of the 20th International Conference on Pattern Recognition, Istanbul, Turkey, 23–26 August 2010; pp. 3121–3124.
65. Sokolova, M.; Lapalme, G. A systematic analysis of performance measures for classification tasks. *Inf. Process. Manag.* **2009**, *45*, 427–437.
66. Chicco, D.; Jurman, G. The advantages of the Matthews correlation coefficient (MCC) over F1 score and accuracy in binary classification evaluation. *BMC Genomics* **2020**, *21*, 6.
67. Platt, J.C. Probabilistic outputs for support vector machines and comparisons to regularized likelihood methods. In *Advances in Large Margin Classifiers*; Smola, A.J., Bartlett, P., Schölkopf, B., Schuurmans, D., Eds.; MIT Press: Cambridge, MA, USA, 1999; pp. 61–74.
68. Guo, C.; Pleiss, G.; Sun, Y.; Weinberger, K.Q. On calibration of modern neural networks. In Proceedings of the 34th International Conference on Machine Learning, Sydney, Australia, 6–11 August 2017; pp. 1321–1330.
69. Vickers, A.J.; Elkin, E.B. Decision curve analysis: A novel method for evaluating prediction models. *Med. Decis. Making* **2006**, *26*, 565–574.
70. Collins, G.S.; Reitsma, J.B.; Altman, D.G.; Moons, K.G.M. Transparent reporting of a multivariable prediction model for individual prognosis or diagnosis (TRIPOD): The TRIPOD statement. *BMJ* **2015**, *350*, g7594.
71. Collins, G.S.; Moons, K.G.M.; Dhiman, P.; Riley, R.D.; Beam, A.L.; Van Calster, B.; Ghassemi, M.; Liu, X.; Reitsma, J.B.; van Smeden, M.; et al. TRIPOD+AI statement: Updated guidance for reporting clinical prediction models that use regression or machine learning methods. *BMJ* **2024**, *385*, e078378.